

# Part 3

## Question 1:

- Implement stochastic gradient and
- compare its performance with that of your gradient descent implementation from Part 1 on the same problem and dataset.
- What happens when both methods are run using the same stepsize?
- Do your observations confirm what has been discussed during the lecture?


We consider the same dataset as in Part 1 and compare the performance of the stochastic descent and gradient descent techniques. To ensure comparability, the following parameters are fixed:

- Number of epochs (eg. number of accesses to the gradient) is set to 30
- $x_0$  the initial estimate is set to all-zeros\$.

As the SG method is not always guaranteed to converge, we opt to compare the methods by making use of a decreasing step size.

- The step parameter is fixed to  $\alpha_k = \frac{0.2}{2^k}$  at iteration  $k$ .

Plot



This run beautifully illustrates typical characteristics of the 2 algorithms when compared one to the other:

As discussed in class, the stochastic gradient algorithm is not a descent method; behaviour which is clearly observed here in that the objective value is sometimes increasing over successive epochs. The oscillatory phase following a phase of quick descent corresponds to class discussion in that we observe the method stalling when the steps get too small to make a big improvement.

Conversely, gradient descent is guaranteed to descend over every successive iteration and we observe this behaviour clearly below.

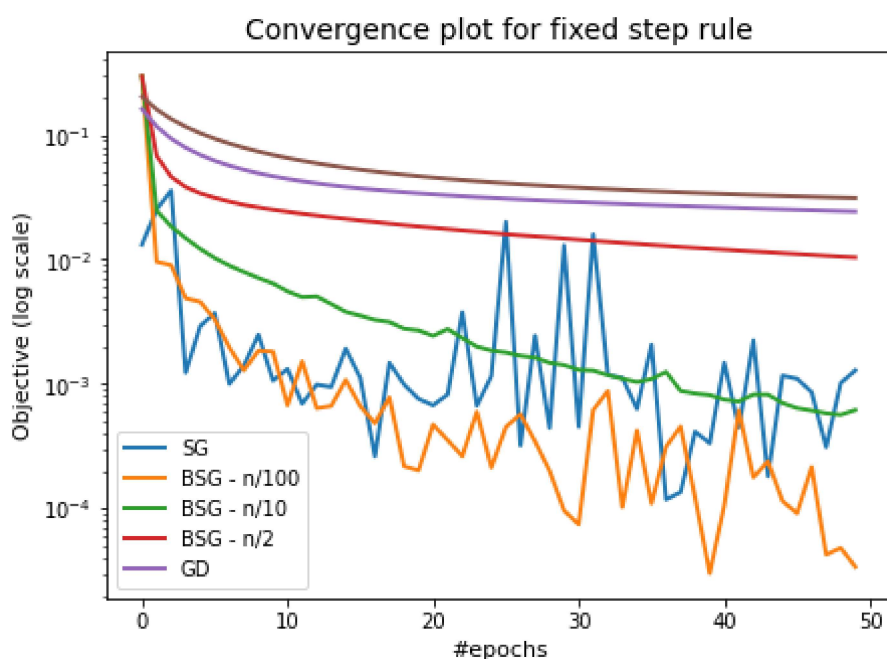
For this method, and for a fixed number of epochs, it seems that SG performs better than the method from which it was derived.

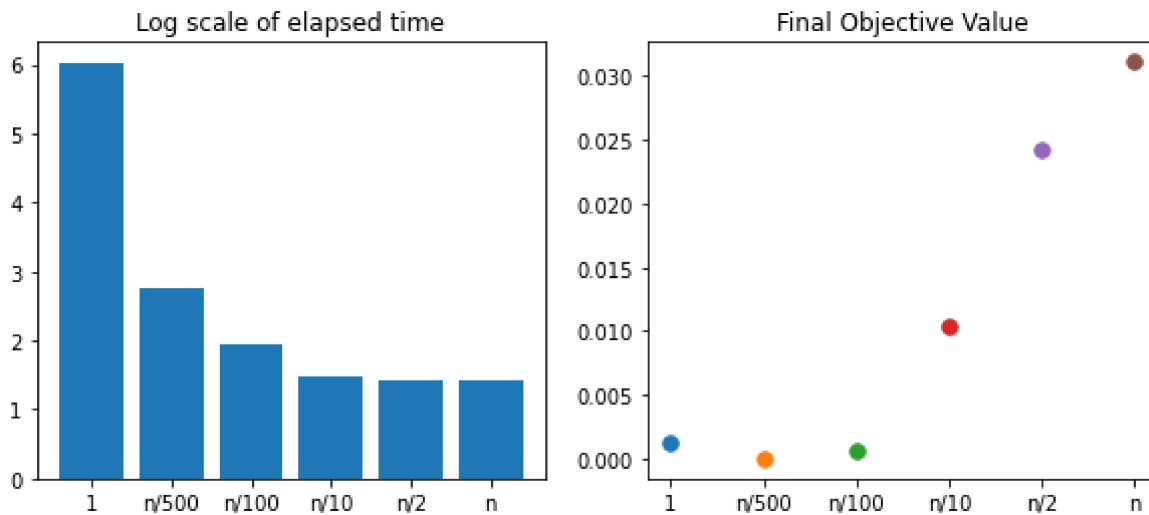
## Question 2:

Find a value for the batch size that realizes a good compromise between gradient descent and stochastic gradient.

We consider several different batch sizes, each for a fixed number of epochs and compare this to the gradient descent method (where the batch size is merely  $n$ , the number of entries).

$$\text{batch} = \left\{1, \frac{n}{100}, \frac{n}{10}, \frac{n}{5}, n\right\}$$





BSG with  $n/100$  seems to perform almost exactly as well as vanilla SG, but with a runtime comparable to smaller batch size.

We conclude that batch stochastic gradient with a batch size of  $n/100$  would be satisfactory as a compromise between SG and GD.

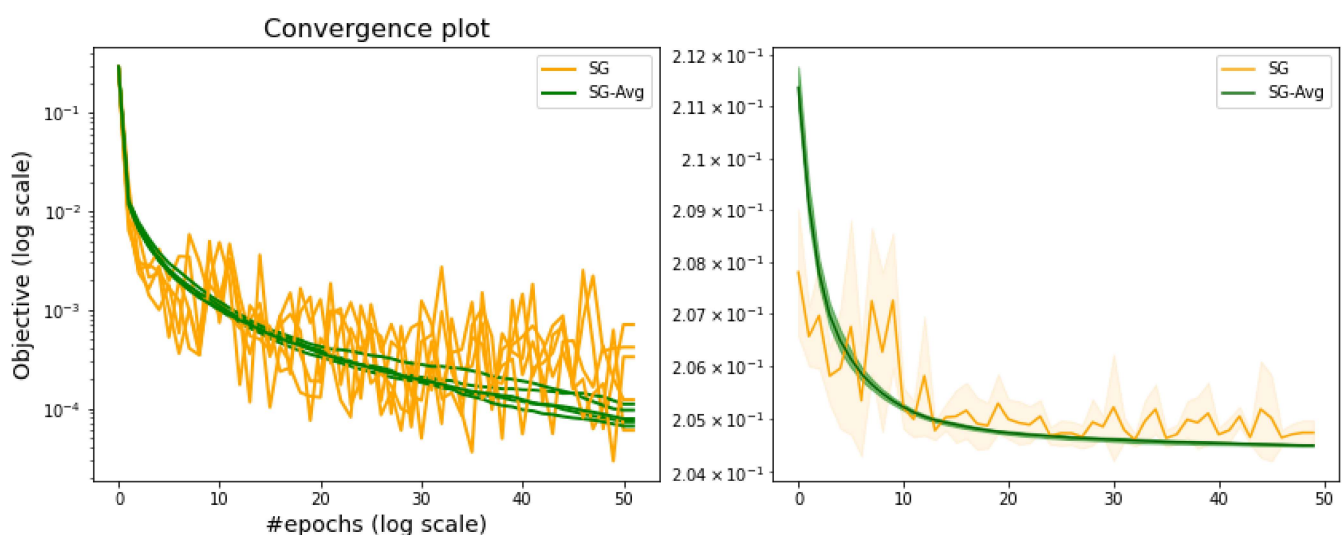
### Question 3:

*Compare your stochastic gradient method with one of the advanced variants seen in class on your selected problem. Discuss your results, and what interpretation you can draw from them.*

We consider SG and compare it with the technique of averaging - a technique designed to reduce the variance. To compare the methods, run 5 repetitions of 50 epochs each for a batch-size fixed to  $\frac{n}{100}$  as determined above and a step size rule as before.

Plotted below on the left is the convergence of each of the 5 runs for SG and SG-averaging respectively.

To the right is the average of each method, with the standard deviation coloured in. We observe that SG-Avg has markedly less variance than that of vanilla SG



As mentioned before, SG is *not* a descent method and so not guaranteed to descend at each iteration. It will, however, descend on average. By averaging over iterates, the SG-Avg method is essentially enforcing the 'descent' behaviour of SGD but still retaining its very good convergence results. Moreover, averaging also decreases the variance in the method.